



# Chamelean

Property-Based Testing for Lean  
via Metaprogramming

Ernest Ng  
(advised by Cody Roux & Mike Hicks)

In collaboration with



# Property-Based Testing

## 1. Write *properties*

Spec for Binary Search Trees (BSTs):


$$\forall x \text{ tree,} \\ \text{isBST tree} \\ \Rightarrow \text{isBST (insert x tree)}$$

# Property-Based Testing (PBT)


## 2. Generate *random inputs*



1, 2, 4, 7, 9, ...



## 3. Test if random inputs satisfy property



# Why should proof assistants support PBT?


Writing proofs is hard & time-consuming

⇒ Use PBT to *test* your specs before starting a proof!

the ability to use the same specifications both for runtime testing and for verification offers potential benefits for both, e.g. for quickly discovering some code and specification errors before embarking on proof,

Banerjee et al. (POPL '25)

# Every major proof assistant has a PBT framework



# Problem: writing PBT generators is hard

Cedar team: writing good generators took 6 months (estimated)



Jane Street devs: writing generators is **"tedious"** & **"high-effort"**

Goldstein et al. (ICSE '24)

# The Constrained Random Generation Problem

Given some proposition  $P$ ,  
*automatically* generate random values satisfying  $P$

# The Constrained Random Generation Problem

Given some *inductive relation*  $P$ ,  
*automatically* generate random values satisfying  $P$



# Inductive Relations

`isBST lo hi tree`

"is `tree` a `BST` containing values between `lo` & `hi`?"

`inductive isBST : Nat → Nat → Tree → Prop where`

`...`

# Inductive Relations

`isBST lo hi tree`

"is `tree` a `BST` containing values between `lo` & `hi`?"

```
inductive isBST : Nat → Nat → Tree → Prop where  
  | BSTLeaf : ∀ lo hi, isBST lo hi Leaf  
  ...
```

# Inductive Relations

`isBST lo hi tree`

"is `tree` a `BST` containing values between `lo` & `hi`?"

```
inductive isBST : Nat → Nat → Tree → Prop where
| BSTLeaf : ∀ lo hi, isBST lo hi Leaf
| BSTNode : ∀ lo hi x l r,
  lo < x < hi →
  isBST lo x l →
  isBST x hi r →
  isBST lo hi (Node x l r)
```

## isBST lo hi tree

"is tree a BST containing values between lo & hi?"

### Function

```
def isBST :  
  Nat → Nat → Tree → Bool :=  
  ...
```

- ✓ Can be executed!
- ✗ Can't encode all properties
- ✗ Hard to reason inductively

### Inductive Relation

```
inductive isBST :  
  Nat → Nat → Tree → Prop where  
  ...
```

- ✓ Facilitates inductive reasoning!
- ✓ Easy to model properties!
- ✗ No computational content  
(Chameleon addresses this issue!)

# Chamelean

1. User specifies a Lean inductive relation

```
inductive isBST :
```


```
  Nat → Nat → Tree → Prop where
```

```
| BSTLeaf : ...
```

```
| BSTNode : ...
```



2. Chamelean derives a generator  
of random data satisfying the relation



(a generator of **Tree**s satisfying **isBST**)

```
#derive_generator (fun tree => isBST lo hi tree)
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  ...
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' => ...
```

pattern-match on **fuel** parameter

(needed to make generators  
structurally decreasing)

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => return Leaf  
  | succ fuel' => ...
```

input **t** unified with **Leaf**  
(talk to me offline about the unification algorithm!)


```
inductive isBST (lo : Nat) (hi : Nat) (t : Tree) : Prop where  
  | BSTLeaf : ∀ lo hi, isBST lo hi Leaf
```



# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => return Leaf  
  | succ fuel' =>  
    backtrack [  
      ...  
    ]
```

**backtrack** keeps picking from a (weighted)  
list of sub-generators until one succeeds



# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => return Leaf  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf), ← input t unified with Leaf  
      ...  
    ]
```

```
inductive isBST (lo : Nat) (hi : Nat) (t : Tree) : Prop where  
  | BSTLeaf : ∀ lo hi, isBST lo hi Leaf
```

# Anatomy of a Generator


```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        ...
```

(Recursive case)

```
| BSTNode : ∀ lo hi x l r,  
  lo < x < hi →  
  isBST lo x l →  
  isBST x hi r →  
  isBST lo hi (Node x l r)
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        let x ← genSuchThat (fun x => lo < x < hi)  
        ...
```



Type class method for the derived generator  
associated with `lo < x < hi`

```
| BSTNode : ∀ lo hi x l r,  
  lo < x < hi →  
  ...
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        let x ← genSuchThat (fun x => lo < x < hi)  
        let l ← genBST lo x fuel'  
        ...  
      )  
    ]
```




Recursively generate left subtree **l**

```
| BSTNode : ∀ lo hi x l r,  
  lo < x < hi →  
  isBST lo x l →  
  ...
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        let x ← genSuchThat (fun x => lo < x < hi)  
        let l ← genBST lo x fuel'  
        let r ← genBST x hi fuel'  
        ...  
      )  
    ]
```



Recursively generate right subtree *r*

```
| BSTNode : ∀ lo hi x l r,  
  lo < x < hi →  
  isBST lo x l →  
  isBST x hi r →  
  ...
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        let x ← genSuchThat (fun x => lo < x < hi)  
        let l ← genBST lo x fuel'  
        let r ← genBST x hi fuel'  
        return (Node x l r))  
    ]
```

Return a **Node**



```
| BSTNode : ∀ lo hi x l r,  
  lo < x < hi →  
  isBST lo x l →  
  isBST x hi r →  
  isBST lo hi (Node x l r)
```

# Anatomy of a Generator

```
def genBST (lo : Nat) (hi : Nat) (fuel : Nat) : Gen Tree :=  
  match fuel with  
  | zero => ...  
  | succ fuel' =>  
    backtrack [  
      (1, return Leaf),  
      (succ fuel', do  
        let x ← genSuchThat (fun x => lo < x < hi)  
        let l ← genBST lo x fuel'  
        let r ← genBST x hi fuel'  
        return (Node x l r)  
      )  
    ]
```

inductive isBST

(lo : Nat)

(hi : Nat)

(t : Tree) : Prop where

| BSTNode :  $\forall$  lo hi x l r,

lo < x < hi  $\rightarrow$

isBST lo x l  $\rightarrow$

isBST x hi r  $\rightarrow$

isBST lo hi (Node x l r)

input **t** unified with (Node x l r)






# Some Generators are Better Than Others

A *naïve* BST generator

`let x ← arbitrary`



Generate some `arbitrary` (unconstrained) `x`

# Some Generators are Better Than Others

A *naïve* BST generator

```
let x ← arbitrary      ← Generate some arbitrary (unconstrained) x
if (lo < x < hi) then  ← Check if lo < x < hi
    ...
    return some (Node x l r)
else
    return none
```

**Bad!**

**For arbitrary  $x$ ,  $\mathbb{P}(\text{lo} < x < \text{hi})$  is low**

# Generator Schedules

A *smarter* BST generator

do

```
let x ← genSuchThat  
      (fun x ⇒ lo < x < hi)  
let l ← genBST lo x fuel'  
let r ← genBST x hi fuel'  
return (Node x l r))
```

- Prioritize constrained generation (`genSuchThat`) over checks
- Rewrite generator based on variable dependencies (i.e. generate `x` before `l` & `r`)

**Testing Theorems, Fully Automatically**

ANONYMOUS AUTHOR(S)

(submitted to POPL '26)

# Chamelean also derives Checkers

inductive Permutation

(l : List Nat)


(l' : List Nat) : Prop



Chamelean derives a function which *checks*

if l, l' are permutations of each other

(this is a semi-decision procedure)



# Chamelean also derives **Enumerators**

`inductive Permutation : List Nat → List Nat → Prop where`

`| Transitivity : ∀ l1 l2 l3,`  
    `Permutation l1 l2 →`  
    `Permutation l2 l3 →`  
    `Permutation l1 l3`

...

$l_2$  doesn't appear in the conclusion of `Transitivity`

⇒ Chamelean ***enumerates***  $l_2$  such that  $l_2$  is also a valid `Permutation`


# Deriving Unconstrained Generators for Inductive Types

inductive Tree

| Leaf : Tree

| Node : Nat → Tree → Tree → Tree

deriving Arbitrary, Enum



# Chamelean implements ideas pioneered in Rocq's QuickChick PBT framework

(talk to me offline about these papers!)

## Testing Theorems, Fully Automatically

ANONYMOUS AUTHOR(S)

Submitted to  
POPL '26



## Computing Correctly with Inductive Relations

Zoe Paraskevopoulou  
z.paraskevopoulou@northeastern.edu  
Northeastern University  
USA

Aaron Eline  
aeline@umd.edu  
University of Maryland, College Park  
USA

Leonidas Lampropoulos  
leonidas@umd.edu  
University of Maryland, College Park  
USA

PLDI '22

## Generating Good Generators for Inductive Relations

LEONIDAS LAMPROPOULOS, University of Pennsylvania, USA  
ZOE PARASKEVOPOULOU, Princeton University, USA  
BENJAMIN C. PIERCE, University of Pennsylvania, USA

POPL '18



# Demo





```

/-- Datatype for binary trees -/
inductive Tree where
| Leaf : Tree
| Node : Nat → Tree → Tree → Tree

/-- `Between lo x hi` means `lo < x < hi` -/
inductive Between : Nat → Nat → Nat → Prop where
| BetweenN : ∀ n m,
  n ≤ m →
  Between n (.succ n) (.succ (.succ m))
| BetweenS : ∀ n m o,
  Between n m o → Between n (.succ m) (.succ o)

#derive_generator (fun (x : Nat) ⇒ Between lo x hi)

/-- `BST lo hi t` describes whether a tree `t` is a BST that
| contains values strictly within `lo` and `hi` -/
inductive BST : Nat → Nat → Tree → Prop where
| BSTLeaf: ∀ lo hi, BST lo hi .Leaf
| BSTNode: ∀ lo hi x l r,
  Between lo x hi →
  BST lo x l →
  BST x hi r →
  BST lo hi (.Node x l r)

#derive_generator (fun (t : Tree) ⇒ BST lo hi t)

```

Derives a generator for  
**x** satisfying **lo < x < hi**

Derives a generator for  
**Trees** satisfying **BST**

# Code for derived generator automatically displayed in VS Code side panel



The image shows the VS Code interface with two panels. The left panel displays a Lean file named `FinalDemo.lean` with the following code:

```
4  /-- Datatype for binary trees -/
5  inductive Tree where
6  | Leaf : Tree
7  | Node : Nat → Tree → Tree → Tree
8
9  /-- `Between lo x hi` means `lo < x < hi` -/
10 inductive Between : Nat → Nat → Nat → Prop where
11 | BetweenN : ∀ n m,
12   n ≤ m →
13   Between n (.succ n) (.succ (.succ m))
14 | BetweenS : ∀ n m o,
15   Between n m o → Between n (.succ m) (.succ o)
16
17 #derive_generator (fun (x : Nat) => Between lo x hi)
18
19 /-- `BST lo hi t` describes whether a tree `t` is a BST that
20 | contains values strictly within `lo` and `hi` -/
21 inductive BST : Nat → Nat → Tree → Prop where
22 | BSTLeaf: ∀ lo hi, BST lo hi .Leaf
23 | BSTNode: ∀ lo hi x l r,
24   Between lo x hi →
25   BST lo x l →
26   BST x hi r →
27   BST lo hi (.Node x l r)
28
29 #derive_generator (fun (t : Tree) => BST lo hi t)
30
```

The right panel shows the `Lean InfoView` for the file, displaying the generated code for the `#derive_generator` directives. The code is as follows:

```
▼ FinalDemo.lean:29:49
▼ Suggestions

Try this generator: instance : ArbitrarySizedSuchThat Tree (fun t_1 => BST lo_1 hi_1 t_1) where
arbitrarySizedST :=
  let rec aux_arb (initSize : Nat) (size : Nat) (lo_1 : Nat) (hi_1 : Nat) : OptionT Plausible.Gen Tree :=
    match size with
    | Nat.zero => OptionTGen.backtrack [(1, return Tree.Leaf)]
    | Nat.succ size' =>
      OptionTGen.backtrack
        [(1, return Tree.Leaf),
         (Nat.succ size', do
           let x <- ArbitrarySizedSuchThat.arbitrarySizedST (fun x => Between lo_1 x hi_1) initSize;
           do
             let l <- aux_arb initSize size' lo_1 x;
             do
               let r <- aux_arb initSize size' x hi_1;
               return Tree.Node x l r)]
         fun size => aux_arb size size lo_1 hi_1

► Messages (1)
► All Messages (2)
```

Clicking on suggested code automatically inserts the generator into your Lean file

```
/-- `BST lo hi t` describes whether a tree `t` is a BST that
| contains values strictly within `lo` and `hi` -/
inductive BST : Nat → Nat → Tree → Prop where
| BSTLeaf: ∀ lo hi, BST lo hi .Leaf
| BSTNode: ∀ lo hi x l r,
  Between lo x hi →
  BST lo x l →
  BST x hi r →
  BST lo hi (.Node x l r)

instance {lo_1 hi_1} : ArbitrarySizedSuchThat Tree (fun t_1 ⇒ BST lo_1 hi_1 t_1) where
  arbitrarySizedST :=
    let rec aux_arb (initSize : Nat) (size : Nat) (lo_1 : Nat) (hi_1 : Nat) : OptionT Plausible.Gen Tree :=
      match size with
      | Nat.zero ⇒ OptionTGen.backtrack [(1, return Tree.Leaf)]
      | Nat.succ size' ⇒
        OptionTGen.backtrack
          [(1, return Tree.Leaf),
           (Nat.succ size', do
             let x ← ArbitrarySizedSuchThat.arbitrarySizedST (fun x ⇒ Between lo_1 x hi_1) initSize;
             do
               let l ← aux_arb initSize size' lo_1 x;
               do
                 let r ← aux_arb initSize size' x hi_1;
                 return Tree.Node x l r)]
    fun size ⇒ aux_arb size size lo_1 hi_1
```

Inserted  
automatically

# Testing a property using Chamelean

Property

$\forall x \text{ lo hi tree,}$   
 $\text{BST lo hi tree} \wedge \text{lo} < x < \text{hi}$   
 $\Rightarrow \text{BST lo hi (insert x tree)}$

Test harness  
(pseudocode)

```
let x ← chooseNat lo hi
let t ← genSuchThat
      (fun tree ⇒ BST lo hi tree)
let t' := insert x t
check (BST lo hi t')
```

} Derived  
generator

← Derived  
checker

(obtained the same  
way as generators)

# Testing a property using Chamelean

```
-- Inserts an element into a tree, respecting the BST invariants -/
def insert (x : Nat) (t : Tree) : Tree :=
  match t with
  | .Leaf ⇒ .Node x .Leaf .Leaf
  | .Node y l r ⇒
    if x < y then
      .Node y (insert x l) r
    else if x > y then
      .Node y l (insert x r)
    else t

-- Test harness for testing the property
`∀ (x : Nat) (t : Tree), BST 0 10 t → BST 0 10 (insert x t)`
for `numTrials` iterations.

-- (Details omitted) -/
def runTests (numTrials : Nat) : IO Unit := ...

-- Uncomment this to run the aforementioned test harness
#eval runTests (numTrials := 10000)
```

## ▼ Messages (1)

### ▼ FinalDemo.lean:82:0

Chamelean: finished 10000 tests, 10000 passed

## ► All Messages (5)

# Falsifying a property using Chamelean

Buggy BST  
insertion function



```
/-- Buggy insertion function: ignores the input tree and
    returns a two-node tree where both values are `x` -/
def buggyInsert (x : Nat) (_ : Tree) : Tree :=
  .Node x (.Node x .Leaf .Leaf) .Leaf

/-- Test harness for testing the property
    `∀ (x : Nat) (t : Tree), BST 0 10 t → BST 0 10 (insert x t)`
    for `numTrials` iterations.

    (Details omitted) -/
def runTests (numTrials : Nat) : IO Unit := ...

-- Uncomment this to run the aforementioned test harness
#eval runTests (numTrials := 10000)
```

▼ Messages (1)

▼ FinalDemo.lean:84:0

Property falsified!

t = Tree.Node 9 (Tree.Node 8 (Tree.Node 7 (Tree.Node 2 (Tree.Leaf)  
(Tree.Leaf)) (Tree.Leaf)) (Tree.Leaf)) (Tree.Leaf)

x = 4

t' = Tree.Node 4 (Tree.Node 4 (Tree.Leaf) (Tree.Leaf)) (Tree.Leaf)

► All Messages (5)

# Case Studies



# Examples

$\Gamma \vdash e : \tau$


Well-typed  
STLC terms



Binary trees  
(BSTs, balanced, complete, ...)

$[^*?@[^*?\.[^*]$

Strings that match  
regular expressions



API calls to a  
Key-Value Store



# Generating Well-Typed STLC Terms

This is an area of active research:

Pałka et al. AST '11  
Fetscher et al. ESOP '15  
Claessen et al. JFP '15  
Midtgaard et al. ICFP '17  
Frank et al. POPL '24

...

Chameleon automatically derives a generator for well-typed terms!

```
λ x : Nat. 1
λ x : Nat. ((3 + x) + 4)
((λ x : Nat. x) 0) + ((λ x : Nat. x + 4) 1)
```

# Aside: the Cedar language

Cutler et al. OOPSLA '24  
Disselkoen et al. FSE '24

DSL developed at AWS  
for rule-based access control:

```
// Example Cedar policy:  
// Interns can't create lists  
forbid(  
    principal in Team::interns,  
    action == CreateList,  
    resource  
);
```

# Ongoing Work: Generating Cedar Terms

$\alpha; \Gamma \vdash \text{true} : \text{True}; \emptyset$	$\alpha; \Gamma \vdash e_1 : E_1; \varepsilon_1$	$s_1 \neq s_2$
$\alpha; \Gamma \vdash \text{false} : \text{False}; \varepsilon$	$\alpha; \Gamma \vdash e_2 : E_2; \varepsilon_2$	$E_1 \neq E_2$
$\alpha; \Gamma \vdash e == e : \text{True}; \emptyset$	$\alpha; \Gamma \vdash e_1 == e_2 : \text{False}; \varepsilon$	$\alpha; \Gamma \vdash E::s_1 == E::s_2 : \text{False}; \varepsilon$
$\alpha; \Gamma \vdash e_1 : \tau_1; \varepsilon_1$	$\tau_1 <: \tau$	$\alpha; \Gamma \vdash e : E_1; \varepsilon$
$\alpha; \Gamma \vdash e_2 : \tau_2; \varepsilon_2$	for some $\tau$	$(E_1 = E) \Rightarrow \tau = \text{True}$
$\alpha; \Gamma \vdash e_1 == e_2 : \text{Bool}; \emptyset$		$(E_1 \neq E) \Rightarrow \tau = \text{False}$
$\alpha; \Gamma \vdash e_1 : E_1; \varepsilon_1$	$\alpha; \Gamma \vdash e_2 : E_2; \varepsilon_2$	
$E_1 \neq E_2$	$M(E_1) = (\_, H)$	$E_2 \notin H$
$\alpha; \Gamma \vdash e_1 \text{ in } e_2 : \text{False}; \varepsilon$	$\alpha; \Gamma \vdash e_1 : E_1; \varepsilon_1$	$\alpha; \Gamma \vdash e_2 : E_2; \varepsilon_2$
$\alpha; \Gamma \vdash e_1 : \text{True}; \varepsilon_1$	$\alpha \cup \varepsilon_1; \Gamma \vdash e_2 : \tau; \varepsilon_2$	$\alpha; \Gamma \vdash e_1 : \text{False}; \varepsilon_1$
$\alpha; \Gamma \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 : \tau; \varepsilon_1 \cup \varepsilon_2$		$\alpha; \Gamma \vdash e_3 : \tau; \varepsilon_3$
$\alpha; \Gamma \vdash e_1 : \text{Bool}; \varepsilon_1$	$\alpha \cup \varepsilon_1; \Gamma \vdash e_2 : \tau; \varepsilon_2$	$\alpha; \Gamma \vdash e_3 : \tau; \varepsilon_3$
$\alpha; \Gamma \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 : \tau; (\varepsilon_1 \cup \varepsilon_2) \cap \varepsilon_3$		$\alpha; \Gamma \vdash e : \tau; \varepsilon$
$\alpha; \Gamma \vdash e : \tau; \varepsilon$	$\alpha; \Gamma \vdash e : \tau; \varepsilon$	$\text{attribute}(f, \tau) = \omega f : \tau_f$ when
$\text{attribute}(f, \tau) = (? f : \tau_f)$	$\text{attribute}(f, \tau) = (\omega f : \tau_f)$	$\tau = \{..., \omega f : \tau_f, ...\}$ , or
$\alpha; \Gamma \vdash e \text{ has } f : \text{Bool}, \{e.f\}$	$\omega = ? \Rightarrow e.f \in \alpha$	$\tau = E$ and
$\alpha; \Gamma \vdash e.f : \tau_f; \emptyset$		$M(E) = (\{..., \omega f : \tau_f, ...\}, \_)$

Lean formalization of Cedar's *static* semantics:


- 29 inductive relations
- Syntax defined via 17 types
- Chamelean can handle 23 / 41 typing rules
- Struggles with typing rules involving complex constraints
  - Algorithm times out! (Talk to me offline for details)

Cutler et al. OOPSLA '24

# Building Chamelean via Metaprogramming

- Before internship: Chamelean prototype produced Lean code via pretty-printed strings
- My job: Implement QuickChick's algorithms using Lean metaprogramming idioms

## Lean Monad Zoo



My prior research: developing PBT tools using OCaml metaprogramming

### **MICA: Automated Differential Testing for OCaml Modules**

ERNEST NG\*, University of Pennsylvania and Cornell University, USA

HARRISON GOLDSTEIN\*, University of Pennsylvania and University of Maryland, USA

BENJAMIN C. PIERCE, University of Pennsylvania, USA

OCaml Workshop '24

# Future / Related Work

# The Golden Age of PBT Research

6 PBT papers to appear at ICFP / OOPSLA '25!

## Bennet: Randomized Specification Testing for Heap-Manipulating Programs

ZAIN K AAMER, University of Pennsylvania, USA  
BENJAMIN C. PIERCE, University of Pennsylvania, USA

## Tuning Random Generators

Property-Based Testing as Probabilistic Programming

RYAN TJOA, University of Washington, USA  
POORVA GARG, University of California, Los Angeles, USA  
HARRISON GOLDSTEIN, University of Maryland, USA  
TODD MILLSTEIN, University of California, Los Angeles, USA  
BENJAMIN C. PIERCE, University of Pennsylvania, USA  
GUY VAN DEN BROECK, University of California, Los Angeles, USA

## We've Got You Covered: Type-Guided Repair of Incomplete Input Generators

PATRICK LAFONTAINE, Purdue University, USA  
ZHE ZHOU, Purdue University, USA  
ASHISH MISHRA, IIT Hyderabad, India  
SURESH JAGANNATHAN, Purdue University, USA  
BENJAMIN DELAWARE, Purdue University, USA

## Teaching Software Specification (Experience Report)

CAMERON MOY, Northeastern University, USA  
DANIEL PATTERSON, Northeastern University, USA

## Lightweight Testing of Persistent Amortized Time Complexity in the Credit Monad

Technical Report, Aug 19, 2025 (v4).

Anton Lorenzen  
University of Edinburgh  
Edinburgh, UK  
anton.lorenzen@ed.ac.uk

☆ [An Empirical Evaluation of Property-Based Testing](#)  
Savitha Ravi, Michael Coblenz

I maintain a PBT bibliography on GitHub:

[ngernest/pbt-bibliography](#)




# Future Work


Automatically produce correctness proofs for derived generators

(ask me offline about how we prove this!)

Soundness: If




is generated, then



is a valid BST

Completeness: All BSTs can be generated



# Future Work

Give users greater control over the distribution of generated values

Allow users to assign probabilities to constructors (à la OCaml QuickCheck)

```
type tree =  
  | Leaf  
  | Node1 of tree * int * tree [@weight 1/2]  
  | Node2 of tree * int * tree [@weight 1/3]  
[@@deriving quickcheck]
```

Tune generators using ideas from the Dice probabilistic language

## Tuning Random Generators

Property-Based Testing as Probabilistic Programming

RYAN TJOA, University of Washington, USA

POORVA GARG, University of California, Los Angeles, USA

HARRISON GOLDSTEIN, University of Maryland, USA

TODD MILLSTEIN, University of California, Los Angeles, USA

BENJAMIN C. PIERCE, University of Pennsylvania, USA

GUY VAN DEN BROECK, University of California, Los Angeles, USA

To appear at OOPSLA '25



# Future Work

Integrate Chamelean with PBT projects by our collaborators

Derive generators using Lean's Aesop proof search tactic

Heuristics for optimizing derived generators

**The Search for Constrained Random Generators**

ANONYMOUS AUTHOR(S)


**Testing Theorems, Fully Automatically**

ANONYMOUS AUTHOR(S)

Both submitted to POPL '26

# Future Work

Benchmark Chamelean's derived generators against their QuickChick counterparts



## ETNA: An Evaluation Platform for Property-Based Testing (Experience Report)

JESSICA SHI, University of Pennsylvania, USA

ALPEREN KELES, University of Maryland, USA

HARRISON GOLDSTEIN, University of Pennsylvania, USA


BENJAMIN C. PIERCE, University of Pennsylvania, USA

LEONIDAS LAMPROPOULOS, University of Maryland, USA

ICFP '23

(extended version submitted to JFP)

# Use Lean's support for live programming to give users greater insight into generated values



Example from Harry Goldstein

# Future Work

Extend Tyche with support for Chamelean  
(VS Code extension for visualizing PBT effectiveness)

## Tyche: Making Sense of Property-Based Testing Effectiveness

Harrison Goldstein

University of Pennsylvania  
Philadelphia, PA, USA  
hgo@seas.upenn.edu

Jeffrey Tao

University of Pennsylvania  
Philadelphia, PA, USA  
jefftao@seas.upenn.edu

Zac Hatfield-Dodds\*

Anthropic  
San Francisco, CA, USA  
zac.hatfield.dodds@gmail.com

Benjamin C. Pierce

University of Pennsylvania  
Philadelphia, PA, USA  
bcpierce@seas.upenn.edu

Andrew Head

University of Pennsylvania  
Philadelphia, PA, USA  
head@seas.upenn.edu

Goldstein et al. UIST '24



# Summary

Chamelean derives the following  
using Lean metaprogramming idioms:

**For inductive relations**

Generators

Enumerators

Checkers

**For algebraic data types**

Generators

Enumerators

# Chamelean is open-source!

ngernest/chamelean



Ongoing work:  
merging Chamelean into Lean's Plausible PBT library

**Feat: Automatically Derive Generators for Algebraic Data Types #35**

 **Open** ngernest wants to merge 60 commits into `leanprover-community:main` from `ngernest:main` 

# Thank you!

Ernest Ng

[ernest@cs.cornell.edu](mailto:ernest@cs.cornell.edu)



# Appendix





# Inductive Relations Can Encode Non-Termination

-- A contrived example: `isZero` holds for `0` and no other `Nat`

`inductive isZero : Nat → Prop where`

`| Zero : isZero 0`

`| NonZero : ∀ n, isZero (succ n) → isZero n`

Example from Paraskevopoulou et al. (PLDI '22), §5

- The derived checker always returns `None` for any non-zero `Nat` (regardless of how much fuel is supplied)
- It will try to satisfy the `isZero (succ n)` hypothesis of the `NonZero` constructor (which is never satisfied!) until it runs out of fuel

# Inductive Relations Can Encode Non-Determinism

Small-step semantics for the untyped  $\lambda$ -calculus,  
extended with non-deterministic choice  $e_1 \sqcap e_2$

$$\frac{}{e_1 \sqcap e_2 \longrightarrow e_1} \text{ (SCHOOSLEFT) }$$

$$\frac{}{e_1 \sqcap e_2 \longrightarrow e_2} \text{ (SCHOOSERIGHT) }$$

Example from CS 6110 Lecture 5

# Functions vs Inductive Relations (Cont.)

Functions	(Inductive) Relations
$f : A \rightarrow B$	$r : A \rightarrow B \rightarrow \text{Prop}$
<b>Total:</b> for every $x : A$ , there is one $f\ x : B$ (termination checker)	<b>Partial:</b> for every $x : A$ , there may be zero or more $y : B$ such that $r\ x\ y$
In Coq: functions are <b>computable</b> <del>terminates : TuringMachine <math>\rightarrow</math> Input <math>\rightarrow</math> bool</del> <del>equal : <math>\mathbb{R} \rightarrow \mathbb{R} \rightarrow</math> bool</del>	May be <b>non-computable</b> terminates : TuringMachine $\rightarrow$ Input $\rightarrow$ Prop equal : $\mathbb{R} \rightarrow \mathbb{R} \rightarrow$ Prop
Proving $f\ x = y$ is <b>automatic</b> if $x, y$ are constants (simpl, reflexivity)	Proving $r\ x\ y$ is <b>manual</b> even if $x, y$ are constants (apply constructors of $r$ )
$\text{eval} : \text{env} \rightarrow \text{stmt} \rightarrow \text{env}$	$\text{big\_step} : \text{env} \rightarrow \text{stmt} \rightarrow \text{env} \rightarrow \text{Prop}$

Taken from CS 6115 (Fall '24) Lecture 7 slides

# Chamelean / QuickChick's Unification Algorithm

$$\begin{aligned}
 \text{unify } u_1 \ u_2 &= \begin{cases} \text{return } () & \text{if } u_1 = u_2 \\ r_1 \leftarrow \kappa[u_1]; r_2 \leftarrow \kappa[u_2]; \text{unifyR } (u_1, r_1) \ (u_2, r_2) & \text{otherwise} \end{cases} \\
 \text{unify } (C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{unifyC } (C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unify } u_1 \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= r_1 \leftarrow \kappa[u_1]; \text{unifyRC } (u_1, r_1) \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unify } (C_1 \ r_{11} \ \dots \ r_{1n}) \ u_2 &= r_2 \leftarrow \kappa[u_2]; \text{unifyRC } (u_2, r_2) \ (C_1 \ r_{11} \ \dots \ r_{1n}) \\
 \\ 
 \text{unifyR } (u_1, \text{undef}_\tau) \ (u_2, r) &= \text{update } u_1 \ u_2 \\
 \text{unifyR } (u_1, r) \ (u_2, \text{undef}_\tau) &= \text{update } u_2 \ u_1 \\
 \text{unifyR } (u_1, u'_1) \ (u_2, r) &= \text{unify } u'_1 \ u_2 \\
 \text{unifyR } (u_1, r) \ (u_2, u'_2) &= \text{unify } u_1 \ u'_2 \\
 \text{unifyR } (-, C_1 \ r_{11} \ \dots \ r_{1n}) \ (-, C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{unifyC } (C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unifyR } (u_1, \text{fixed}) \ (u_2, \text{fixed}) &= \text{equality } u_1 \ u_2; \text{update } u_1 \ u_2 \\
 \text{unifyR } (u_1, \text{fixed}) \ (-, C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{match } u_1 \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unifyR } (-, C_1 \ r_{11} \ \dots \ r_{1n}) \ (u_2, \text{fixed}) &= \text{match } u_2 \ (C_1 \ r_{11} \ \dots \ r_{1n}) \\
 \\ 
 \text{unifyC } (C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= \begin{cases} \text{fold } \text{unify } (\overline{r_{1i}, r_{2i}}) & \text{if } C_1 = C_2 \text{ and } n = m \\ \perp & \text{otherwise} \end{cases} \\
 \\ 
 \text{unifyRC } (u, \text{undef}_\tau) \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{update } u_1 \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unifyRC } (u, u') \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= r \leftarrow \kappa[u']; \text{unifyRC } (u', r) \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unifyRC } (u, \text{fixed}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{match } u \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \text{unifyRC } (u, C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) &= \text{unifyC } (C_1 \ r_{11} \ \dots \ r_{1n}) \ (C_2 \ r_{21} \ \dots \ r_{2m}) \\
 \\ 
 \text{match } u \ (C \ r_1 \ \dots \ r_n) &= \bar{p} \leftarrow \text{mapM } \text{matchAux } \bar{r}; \text{pattern } u \ (C \ \bar{p}) \\
 \\ 
 \text{matchAux } (C \ \bar{r}) &= \bar{p} \leftarrow \text{mapM } \text{matchAux } \bar{r}; \text{return } (C \ \bar{p}) \\
 \text{matchAux } u &= r \leftarrow \kappa[u]; \text{case } r \text{ of} \quad \begin{aligned} &\text{undef}_\tau \Rightarrow \text{update } u \ \text{fixed} \\ &| \text{fixed} \Rightarrow u' \leftarrow \text{fresh}; \text{equality } u' \ u; \text{update } u' \ u; \text{return } u' \\ &| u' \Rightarrow \text{matchAux } u' \\ &| C \ \bar{r} \Rightarrow \bar{p} \leftarrow \text{mapM } \text{matchAux } \bar{r}; \text{return } (C \ \bar{p}) \end{aligned}
 \end{aligned}$$

Fig. 3. Unification algorithm

- Works on *unknowns* (set of values that variables can take on during generation)
  - Difference: unknowns can be provided as *inputs* to generators (i.e. they have one single fixed value, but that value is unknown at compile time)
- Implemented using the State & Option monads!

## FAQ

**Q:** Why can't we just use Rocq/Lean's internal unifier?

**A:** Rocq's unifier works on Rocq terms, whereas we need to unify our representation of knowledge of the state of each unknown at any particular point

Lampropoulos et al. POPL '18

# General Structure of a Generator

```
for each  $(u, C \bar{p})$  in patterns:  
  match u with  
  |  $C \bar{p} \Rightarrow \dots$     for each  $(u_1, u_2)$  in equalities:  
    if  $u_1 = u_2$  then ...  
      for each constraint (e.g.,  $x < hi$ ):  
        do!  $x \leftarrow \text{arbST} \dots$  (* instantiations *)  
        if  $(x < hi)?$  then  $\dots$  (* checks *)  
          Final result:  
          ret (Some  $\dots$ )  
        else ret None  
      else ret None  
  | _  $\Rightarrow$  ret None
```

Fig. 1. General Structure of each sub-generator

Lampropoulos et al. POPL '18

# Checker Correctness Proof (Sketch)

## Soundness

$\forall \text{ fuel},$   
 $\text{check } (P \ e_1 \ \dots \ e_m) \ \text{fuel} = \text{Some true}$   
 $\Rightarrow P \ e_1 \ \dots \ e_m$

By induction on the `fuel` argument

## Completeness

$P \ e_1 \ \dots \ e_m$   
 $\Rightarrow \exists \text{ fuel},$   
 $\text{check } (P \ e_1 \ \dots \ e_m) \ \text{fuel} = \text{Some true}$

By induction on the derivation of  $P \ e_1 \ \dots \ e_m$

Proof for generators is more complex (requires reasoning about generators' **support**)

See §5 of Paraskevopoulou et al. (PLDI '22) for more details!

# Checkers Invoking Enumerators: An Example

$$\frac{\Gamma \vdash e_1 : \tau_1 \rightarrow \tau_2 \quad \Gamma \vdash e_2 : \tau_1}{\Gamma \vdash e_1 e_2 : \tau_2} \text{ TApp}$$

```
-- `typing Γ e τ` is the typing judgment `Γ ⊢ e : τ`  
inductive typing: context → term → type → Prop where  
  | TApp : ∀ Γ e1 e2 τ1 τ2,  
    typing Γ e2 τ1 →  
    typing Γ e1 (Fun τ1 τ2) →  
    typing Γ (App e1 e2) τ2  
  ...
```

**τ1** doesn't appear in the conclusion of **TApp**, but is used in hypotheses  
⇒ the derived checker for **typing** needs to enumerate **τ1**

# Checkers Invoking Enumerators: An Example

```
| TApp :  $\forall \Gamma \ e1 \ e2 \ \tau1 \ \tau2,$   
    typing  $\Gamma \ e2 \ \tau1 \rightarrow$   
    typing  $\Gamma \ e1 \ (\text{Fun } \tau1 \ \tau2) \rightarrow$   
    typing  $\Gamma \ (\text{App } e1 \ e2) \ \tau2$ 
```

```
match e,  $\tau$  with  
| App e1 e2,  $\tau2 \Rightarrow$  do  
    let  $\tau1 \leftarrow$  enumerateSuchThat  
        (fun  $\tau \Rightarrow$  typing  $\Gamma \ e2 \ \tau$ )  
    check (typing  $\Gamma \ e1 \ (\text{Fun } \tau1 \ \tau2)$ )  
| ...
```

$\tau1$  doesn't appear in the conclusion of TApp, but is used in hypotheses  
 $\Rightarrow$  the derived checker for typing needs to enumerate  $\tau1$



Given some inductive relation  $P$ , Chamelean derives:

## Generators

(random)

`Gen (Option  $\alpha$ )`

Produce  $\alpha$ 's  
satisfying  $P$

## Enumerators

(deterministic)

`List (Option  $\alpha$ )`

(produced lazily)

## Checkers

(semi-decision procedures)

`Option Bool`

(None = out of fuel)

Checks if  
 $P(\alpha)$  holds

# Smarter constraint ordering

```
inductive BST : Nat → Nat → Tree → Prop where
| BSTNode: ∀ lo hi x l r,
    lo < x < hi →
    BST lo x l →
    BST x hi r →
    BST lo hi (Node x l r)
```

Behavior of derived generator :

1. Generate  $x$  such that  $lo < x < hi$
2. Generate left & right subtrees

# Smarter constraint ordering

inductive BST : Nat → Nat → Tree → Prop where

| BSTNode:  $\forall$  lo hi x l r,

BST lo x l →

BST x hi r →

**lo < x < hi →**

-- Order of hypotheses swapped

BST lo hi (Node x l r)

1. Generate some unconstrained x
2. Generate left & right subtrees
3. Check that lo < x < hi

This check will often fail!

(For arbitrarily generated x,  $\text{Pr}[\text{lo} < x < \text{hi}]$  is low)